

# TEACHNOOK

## MAJOR PROJECT

(FEB – BATCH)

MADE BY GROUP– DS-02-SPB3

## EXTERNAL DATA ANALYSIS

### **Dataset: -**

#### **PROBLEM STATEMENT: -**

Q - In this kernel, we try to make some predictions where we have to determine whether a person makes an income over 50K in a year. We implement Random Forest Classification with Python and Scikit-Learn. So, to answer this question, we had built some Random Forest classifier to predict whether a person can make their income over 50K in a Year or not?

```
✓ 1s [29] import pandas as pd
import numpy as np
import seaborn as sns
import sklearn as sk
import matplotlib.pyplot as plt
```

```
✓ 0s [2] path = "/content/drive/MyDrive/data_csv.csv"
```

```
✓ 0s ▶ df = pd.read_csv(path)
print(df)
```

```
↗
```

	age	workclass	fnlwgt	education	education-num \	
0	39	state-gov	77526	bachelors	13	
1	50	self-emp-not-inc	83311	bachelors	13	
2	38	private	215646	hs-grad	9	
3	53	private	234721	11th	7	
4	28	private	338409	bachelors	13	

	marital-status	occupation	relationship	race	sex
0	never-married	adm-cierical	not-in-family	white	male
1	married-civ-spouse	exec-managerial	husband	white	male
2	divorced	handlers-cleaners	not-in-family	white	male
3	married-civ-spouse	handlers-cleaners	husband	black	male
4	married-civ-spouse	prof-specially	wife	black	female

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	US	<=50k
1	0	0	13	US	<=50k
2	0	0	40	US	<=50k
3	0	0	40	US	<=50k
4	0	0	40	Cuba	<=50k

## GOING THROUGH BASIC INFORMATION:-

#1 Print the 'Workclass' column from the dataset.

```
✓ 0s ▶ #print workclass column form the dataset

print(df['workclass'].unique())
```

```
['state-gov' 'self-emp-not-inc' 'private']
```

#2 Show the dimensions of the given data set.



0s



```
#view the dimensions of the data
print('The shape of the dataset : ', df.shape)
```

The shape of the dataset : (5, 15)

#3 show the preview the dataset.



0s



```
#Preview the dataset
```

```
df.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	income
0	39	state-gov	77526	bachelors	13	never-married	adm-clerical	not-in-family	white	male	2174	0	40	US	<=50k
1	50	self-emp-not-inc	83311	bachelors	13	married-civ-spouse	exec-managerial	husband	white	male	0	0	13	US	<=50k
2	38	private	215646	hs-grad	9	divorced	handlers-cleaners	not-in-family	white	male	0	0	40	US	<=50k
3	53	private	234721	11th	7	married-civ-spouse	handlers-cleaners	husband	black	male	0	0	40	US	<=50k
4	28	private	338409	bachelors	13	married-civ-spouse	prof-specialty	wife	black	female	0	0	40	Cuba	<=50k



#### #4 Renaming columns.

✓  
0s



```
#Rename the column names
```

```
col_names = ['age', 'workclass', 'fnlwgt', 'education', 'education_num',  
             'marital_status', 'occupation', 'relationship',  
             'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week',  
             'native_country', 'income']
```

```
df.columns = col_names
```

```
df.columns
```

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',  
      'marital_status', 'occupation', 'relationship', 'race', 'sex',  
      'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',  
      'income'],  
      dtype='object')
```

#### #5 Show the summery of the dataset.

✓  
0s



```
#View summery
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5 entries, 0 to 4  
Data columns (total 15 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   age                   5 non-null     int64  
1   workclass              5 non-null     object  
2   fnlwgt                 5 non-null     int64  
3   education              5 non-null     object  
4   education_num          5 non-null     int64  
5   marital_status         5 non-null     object  
6   occupation              5 non-null     object  
7   relationship            5 non-null     object  
8   race                   5 non-null     object  
9   sex                    5 non-null     object  
10  capital_gain            5 non-null     int64  
11  capital_loss            5 non-null     int64  
12  hours_per_week          5 non-null     int64  
13  native_country          5 non-null     object  
14  income                  5 non-null     object  
dtypes: int64(6), object(9)  
memory usage: 728.0+ bytes
```

## #6 Check the datatype of each column in the dataset.

✓  
0s



```
#Check the data type of the column  
df.dtypes
```

```
age          int64  
workclass    object  
fnlwgt       int64  
education    object  
education_num int64  
marital_status object  
occupation   object  
relationship object  
race         object  
sex          object  
capital_gain int64  
capital_loss int64  
hours_per_week int64  
native_country object  
income       object  
dtype: object
```

## #7 Show the statistical properties of the dataset.

✓  
0s



```
#view statistical properties of the data set  
df.describe()
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
count	5.00000	5.000000	5.000000	5.000000	5.0	5.000000
mean	41.60000	189922.600000	11.000000	434.800000	0.0	34.600000
std	10.06479	110358.314006	2.828427	972.242357	0.0	12.074767
min	28.00000	77526.000000	7.000000	0.000000	0.0	13.000000
25%	38.00000	83311.000000	9.000000	0.000000	0.0	40.000000
50%	39.00000	215646.000000	13.000000	0.000000	0.0	40.000000
75%	50.00000	234721.000000	13.000000	0.000000	0.0	40.000000
max	53.00000	338409.000000	13.000000	2174.000000	0.0	40.000000

#8 Check if there is any missing value in the given dataset.



0s



```
#Check the missing value
df.isnull().sum()
```

```
age          0
workclass    0
fnlwgt       0
education    0
education_num 0
marital_status 0
occupation   0
relationship 0
race         0
sex          0
capital_gain 0
capital_loss 0
hours_per_week 0
native_country 0
income       0
dtype: int64
```

#9 Find the categorical variable.



0s



```
#Find categorical variable
categorical = [var for var in df.columns if df[var].dtype=='O']

print('There are {} categorical variables\n'.format(len(categorical)))


print('The categorical variables are :\n\n', categorical)
```

There are 9 categorical variables

The categorical variables are :


```
['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country', 'income']
```

## #10 Review the categorical variable.

✓ 0s  #Priview categorical variable  
df[categorical].head()


	workclass	education	marital_status	occupation	relationship	race	sex	native_country	income
0	state-gov	bachelors	never-married	adm-cierical	not-in-family	white	male	US	<=50k
1	self-emp-not-inc	bachelors	married-civ-spouse	exec-managerial	husband	white	male	US	<=50k
2	private	hs-grad	divorced	handlers-cleaners	not-in-family	white	male	US	<=50k
3	private	11th	married-civ-spouse	handlers-cleaners	husband	black	male	US	<=50k
4	private	bachelors	married-civ-spouse	prof-specially	wife	black	female	Cuba	<=50k

## #11 Check the missing value from 'income' column.

✓ 0s  #Exploring the variables  
#Explore income target variable  
# check for missing values  
  
df['income'].isnull().sum()



## #12 show any unique number from the column 'Income'.

✓ 0s  #We can see that there are no missing values in the income target variable.  
  
# view number of unique values  
  
df['income'].nunique()

1

#13 Check the unique value from the column 'income'.

✓  
0s



```
#There are 1 unique values in the income variable.  
  
# view the unique values  
  
df['income'].unique()  
  
array(['<=50k'], dtype=object)
```

#14 show the percentage of frequency distribution of values.

✓  
0s



```
# view percentage of frequency distribution of values  
  
df['income'].value_counts()/len(df)  
  
<=50k    1.0  
Name: income, dtype: float64
```

#15 Count the number of rows in the given dataset.



```
#show number of rows in a given dataset  
len(df)
```

5



#16 Show the minimum 'age' from the given dataset

✓  
0s



```
#Show the minimum 'age'  
print(min(df['age']))
```

28

#17 Show the maximum 'age' from the given dataset.

✓  
0s



```
#Show the maximum 'age'  
print(max(df['age']))
```



53

#18 See if there is any 'capital\_gain' in the given data set.

✓  
0s



```
# See if there is any 'capital_gain'.  
print(df['capital_gain'])
```

```
0      2174  
1         0  
2         0  
3         0  
4         0  
Name: capital_gain, dtype: int64
```

#19 See if there is any 'capital\_loss' in the given data set.

✓  
0s



```
#See if there is any 'capital_loss'.  
print(df['capital_loss'])
```

```
0      0  
1      0  
2      0  
3      0  
4      0  
Name: capital_loss, dtype: int64
```

#20 See the Maximum 'capital\_gain'.

✓  
0s



```
#See the Maximum 'capital_gain'.  
print(df['capital_gain'].max())
```

```
2174
```

#21 See the Minimum 'capital\_loss'.

✓  
0s




```
#See the minimum 'capital_loss'.  
print(df['capital_gain'].min())
```

```
0
```

#22 Show which 'workclass' has maximum 'capital\_gain'

✓  
0s



```
#The is 'capital_gain' in the given data set with the 'workclass'.
```

```
print(df[['workclass','capital_gain']].max())
```

```
workclass      state-gov  
capital_gain    2174  
dtype: object
```

## Predictions: -

### 1. The problem statement

In this kernel, We try to make predictions where the prediction task is to determine whether a person makes over 50K a year. We implement Random Forest Classification with Python and Scikit-Learn. So, to answer the question, We build a Random Forest classifier to predict whether a person makes over 50K a year or not.

### 2. Importing the libraries

```
In [124]: import numpy as np
import pandas as pd
```

```
In [125]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

sns.set(style="whitegrid")
```

```
In [126]: import warnings

warnings.filterwarnings('ignore')
```

### 4. Exploratory data analysis (EDA)

```
In [28]: # print the shape
print('The shape of the dataset : ', df.shape)
```

The shape of the dataset : (32561, 15)

We can see that there are 32561 instances and 15 attributes in the data set.

```
In [29]: df.head()
```

```
Out[29]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	income
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

### 3. Importing the dataset

```
In [127]: path = r'file:///C:/Users/hp/Downloads/income_evaluation.csv'

df = pd.read_csv(path)
```

### 4. Exploratory data analysis (EDA)

```
In [28]: # print the shape
print('The shape of the dataset : ', df.shape)
```

The shape of the dataset : (32561, 15)

We can see that there are 32561 instances and 15 attributes in the data set.

```
In [29]: df.head()
```

```
Out[29]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	income
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-	Not-in-family	White	Male	0	0	40	United-	<=50K

```
In [129]: col_names = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status', 'occupation', 'relationship',
                      'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country', 'income']

df.columns = col_names

df.columns
```

```
Out[129]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
                 'marital_status', 'occupation', 'relationship', 'race', 'sex',
                 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
                 'income'],
                dtype='object')
```

```
In [130]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education_num         32561 non-null  int64
5   marital_status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital_gain          32561 non-null  int64
```

```
In [129]: col_names = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status', 'occupation', 'relationship',
                      'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country', 'income']

df.columns = col_names

df.columns
```

```
Out[129]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
                 'marital_status', 'occupation', 'relationship', 'race', 'sex',
                 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
                 'income'],
                dtype='object')
```

```
In [130]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education_num         32561 non-null  int64
5   marital_status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital_gain          32561 non-null  int64
```

## FINDINGS: -

### Findings

- We can see that the dataset contains 9 character variables and 6 numerical variables.
- `income` is the target variable.
- There are no missing values in the dataset. I will explore this later,

```
In [131]: df.dtypes
```

```
Out[131]: age                int64
workclass                object
fnlwgt                  int64
education                object
education_num            int64
marital_status           object
occupation               object
relationship             object
race                    object
sex                     object
capital_gain            int64
capital_loss            int64
hours_per_week          int64
native_country           object
income                  object
dtype: object
```

```
In [134]: df.describe(include='all')
```

```
Out[134]:
```

```
In [130]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education_num         32561 non-null  int64
5   marital_status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital_gain          32561 non-null  int64
11  capital_loss          32561 non-null  int64
12  hours_per_week        32561 non-null  int64
13  native_country        32561 non-null  object
14  income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```



```
In [36]: # check for missing values
```

```
df.isnull().sum()
```

```
Out[36]: age                0
workclass                0
fnlwgt                  0
education                0
education_num           0
marital_status          0
occupation              0
relationship            0
race                    0
sex                     0
capital_gain            0
capital_loss            0
hours_per_week          0
native_country          0
income                  0
dtype: int64
```

```
In [134]: df.describe(include='all')
```

```
Out[134]:
```

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss
count	32561.000000	32561	3.256100e+04	32561	32561.000000	32561	32561	32561	32561	32561	32561.000000	32561.000000
unique	NaN	9	NaN	16	NaN	7	15	6	5	2	NaN	NaN
top	NaN	Private	NaN	HS-grad	NaN	Married-civ-spouse	Prof-specialty	Husband	White	Male	NaN	NaN
freq	NaN	22696	NaN	10501	NaN	14976	4140	13193	27816	21790	NaN	NaN
mean	38.581647	NaN	1.897784e+05	NaN	10.080679	NaN	NaN	NaN	NaN	NaN	1077.648844	87.303830
std	13.640433	NaN	1.055500e+05	NaN	2.572720	NaN	NaN	NaN	NaN	NaN	7385.292085	402.960219
min	17.000000	NaN	1.228500e+04	NaN	1.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000
25%	28.000000	NaN	1.178270e+05	NaN	9.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000
50%	37.000000	NaN	1.783560e+05	NaN	10.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000
75%	48.000000	NaN	2.370510e+05	NaN	12.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000
max	90.000000	NaN	1.484705e+06	NaN	16.000000	NaN	NaN	NaN	NaN	NaN	99999.000000	4356.000000

```
In [38]: def initial_eda(df):
    if isinstance(df, pd.DataFrame):
        total_na = df.isna().sum().sum()
        print("Dimensions : %d rows, %d columns" % (df.shape[0], df.shape[1]))
        print("Total NA Values : %d " % (total_na))
        print("%38s %10s    %10s %10s" % ("Column Name", "Data Type", "#Distinct", "NA Values"))
        col_name = df.columns
        dtyp = df.dtypes
        uniq = df.nunique()
        na_val = df.isna().sum()
        for i in range(len(df.columns)):
            print("%38s %10s    %10s %10s" % (col_name[i], dtyp[i], uniq[i], na_val[i]))

    else:
        print("Expect a DataFrame but got a %15s" % (type(df)))
```

```
In [39]: initial_eda(df)
```

```
Dimensions : 32561 rows, 15 columns
Total NA Values : 0
```

Column Name	Data Type	#Distinct	NA Values
age	int64	73	0
workclass	object	9	0
fnlwgt	int64	21648	0
education	object	16	0
education_num	int64	16	0
marital_status	object	7	0
occupation	object	15	0
relationship	object	6	0
race	object	5	0

## INTERPRETATION: -

### Interpretation

We can see that there are no missing values in the dataset.

```
In [37]: #assert that there are no missing values in the dataframe

assert pd.notnull(df).all().all()
```

### Interpretation

- The above command does not throw any error. Hence, it is confirmed that there are no missing or negative values in the dataset.
- All the values are greater than or equal to zero excluding character values.

```
In [38]: def initial_eda(df):
    if isinstance(df, pd.DataFrame):
        total_na = df.isna().sum().sum()
        print("Dimensions : %d rows, %d columns" % (df.shape[0], df.shape[1]))
        print("Total NA Values : %d " % (total_na))
        print("%38s %10s    %10s %10s" % ("Column Name", "Data Type", "#Distinct", "NA Values"))
        col_name = df.columns
        dtyp = df.dtypes
        uniq = df.nunique()
        na_val = df.isna().sum()
        for i in range(len(df.columns)):
            print("%38s %10s    %10s %10s" % (col_name[i], dtyp[i], uniq[i], na_val[i]))

    else:
        print("Expect a DataFrame but got a %15s" % (type(df)))
```



```
In [39]: initial_eda(df)
```

```
Dimensions : 32561 rows, 15 columns
```

```
Total NA Values : 0
```

Column Name	Data Type	#Distinct	NA Values
age	int64	73	0
workclass	object	9	0
fnlwgt	int64	21648	0
education	object	16	0
education_num	int64	16	0
marital_status	object	7	0
occupation	object	15	0
relationship	object	6	0
race	object	5	0
sex	object	2	0
capital_gain	int64	119	0
capital_loss	int64	92	0
hours_per_week	int64	94	0
native_country	object	42	0
income	object	2	0

## EXPLORE CATEGORICAL VARIABLES: -

### 4.1. Explore Categorical Variables

```
In [135]: categorical = [var for var in df.columns if df[var].dtype=='O']
```

```
print(f"There are {len(categorical)} categorical variables\n")
```

```
print('The categorical variables are :\n\n', categorical)
```

```
There are 9 categorical variables
```

```
The categorical variables are :
```

```
['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country', 'income']
```

### 4.2 Preview categorical variables

```
In [41]: df[categorical].head()
```

```
Out[41]:
```

	workclass	education	marital_status	occupation	relationship	race	sex	native_country	income
0	State-gov	Bachelors	Never-married	Adm-clerical	Not-in-family	White	Male	United-States	<=50K
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States	<=50K
2	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States	<=50K
3	Private	11th	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	United-States	<=50K
4	Private	Bachelors	Married-civ-spouse	Prof-specialty	Wife	Black	Female	Cuba	<=50K

```
In [39]: initial_eda(df)
```

Dimensions : 32561 rows, 15 columns

Total NA Values : 0

Column Name	Data Type	#Distinct	NA Values
age	int64	73	0
workclass	object	9	0
fnlwgt	int64	21648	0
education	object	16	0
education_num	int64	16	0
marital_status	object	7	0
occupation	object	15	0
relationship	object	6	0
race	object	5	0
sex	object	2	0
capital_gain	int64	119	0
capital_loss	int64	92	0
hours_per_week	int64	94	0
native_country	object	42	0
income	object	2	0

There are 2 unique values in the `income` variable.

```
In [46]: # view the unique values
```

```
df['income'].unique()
```

```
Out[46]: array([' <=50K', ' >50K'], dtype=object)
```

The two unique values are `<=50K` and `>50K`.

```
In [47]: # view the frequency distribution of values
```

```
df['income'].value_counts()
```

```
Out[47]: <=50K    24720
>50K        7841
Name: income, dtype: int64
```

```
In [136]: # visualize frequency distribution of income variable
```

```
f,ax=plt.subplots(1,2,figsize=(18,8))
```

```
ax[0] = df['income'].value_counts().plot.pie(explode=[0,0],autopct='%1.1f%%',ax=ax[0],shadow=True)
ax[0].set_title('Income Share')
```

```
#f, ax = plt.subplots(figsize=(6, 8))
```

```
ax[1] = sns.countplot(x="income", data=df, palette="ocean_r")
ax[1].set_title("Frequency distribution of income variable")
```

### 4.3 Summary of categorical variables

- There are 9 categorical variables in the dataset.
- The categorical variables are given by `workclass`, `education`, `marital_status`, `occupation`, `relationship`, `race`, `sex`, `native_country` and `income`.
- `income` is the target variable.

### 4.4 Explore the variables

Explore `income` target variable

```
In [44]: # check for missing values
df['income'].isnull().sum()
```


Out[44]: 0

We can see that there are no missing values in the `income` target variable.

```
In [45]: # view number of unique values
df['income'].nunique()
```

Out[45]: 2

```
plt.show()
```



```
In [47]: # view the frequency distribution of values
df['income'].value_counts()
```

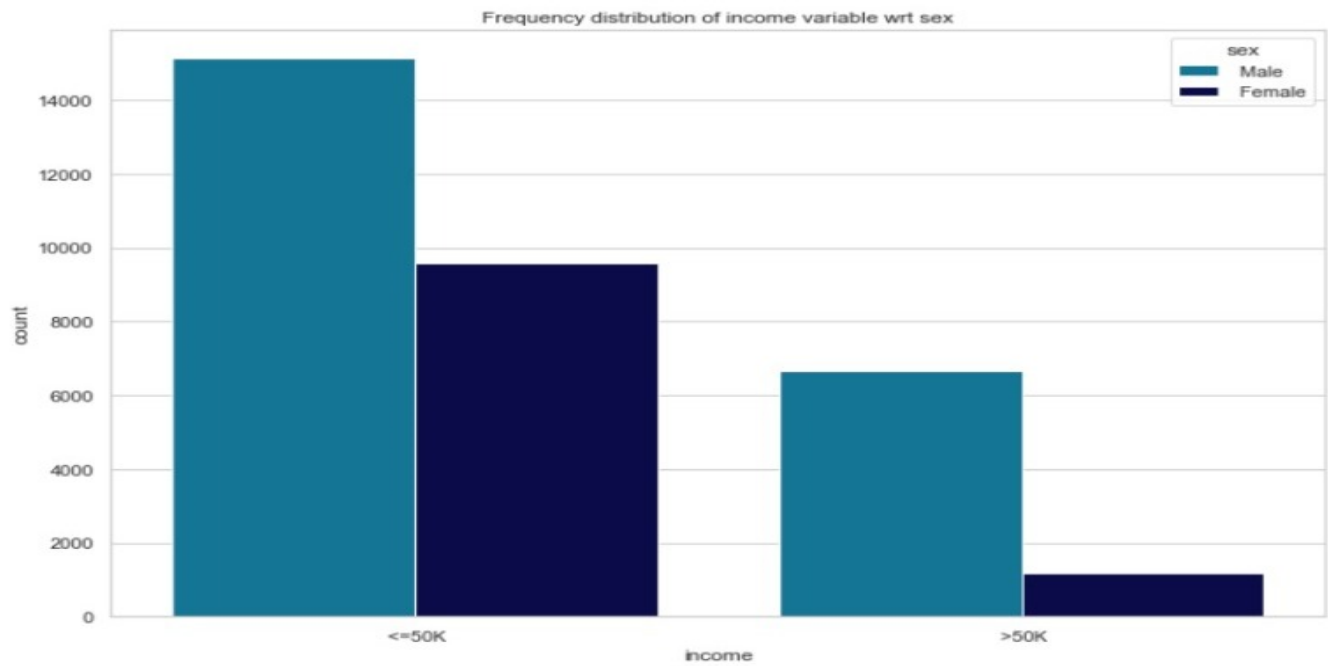
```
Out[47]: <=50K    24720
         >50K     7841
         Name: income, dtype: int64
```

```
In [136]: # visualize frequency distribution of income variable
f,ax=plt.subplots(1,2,figsize=(18,8))

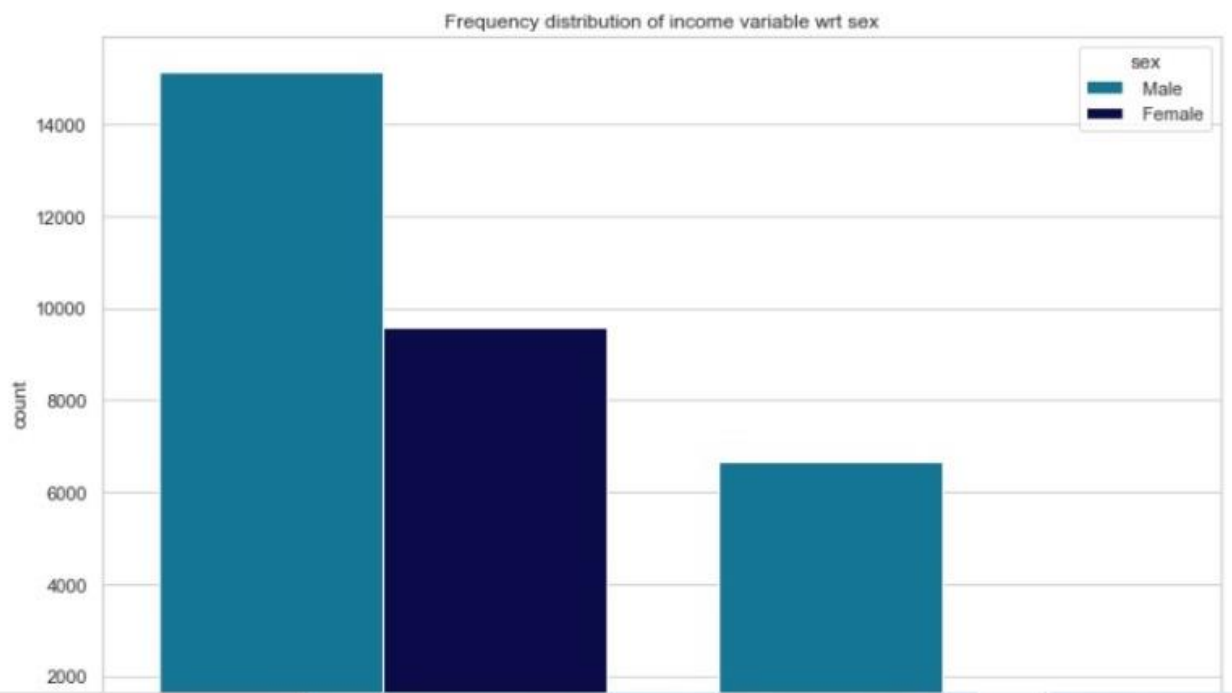
ax[0] = df['income'].value_counts().plot.pie(explode=[0,0],autopct='%1.1f%%',ax=ax[0],shadow=True)
ax[0].set_title('Income Share')

#f, ax = plt.subplots(figsize=(6, 8))
ax[1] = sns.countplot(x="income", data=df, palette="ocean_r")
ax[1].set_title("Frequency distribution of income variable")

plt.show()
```



```
In [143]: categorical_features=["sex","race"]
for feature in categorical_features:
    f,ax = plt.subplots(figsize=(12, 8))
    ax = sns.countplot(x="income", hue=feature, data=df, palette="ocean_r")
    ax.set_title(f"Frequency distribution of income variable wrt {feature}")
    plt.show()
```



### Interpretation

- We can see that whites make more money than non-whites in both the income categories.

### Explore `workclass` variable

```
In [53]: # check number of unique labels
df.workclass.nunique()
```

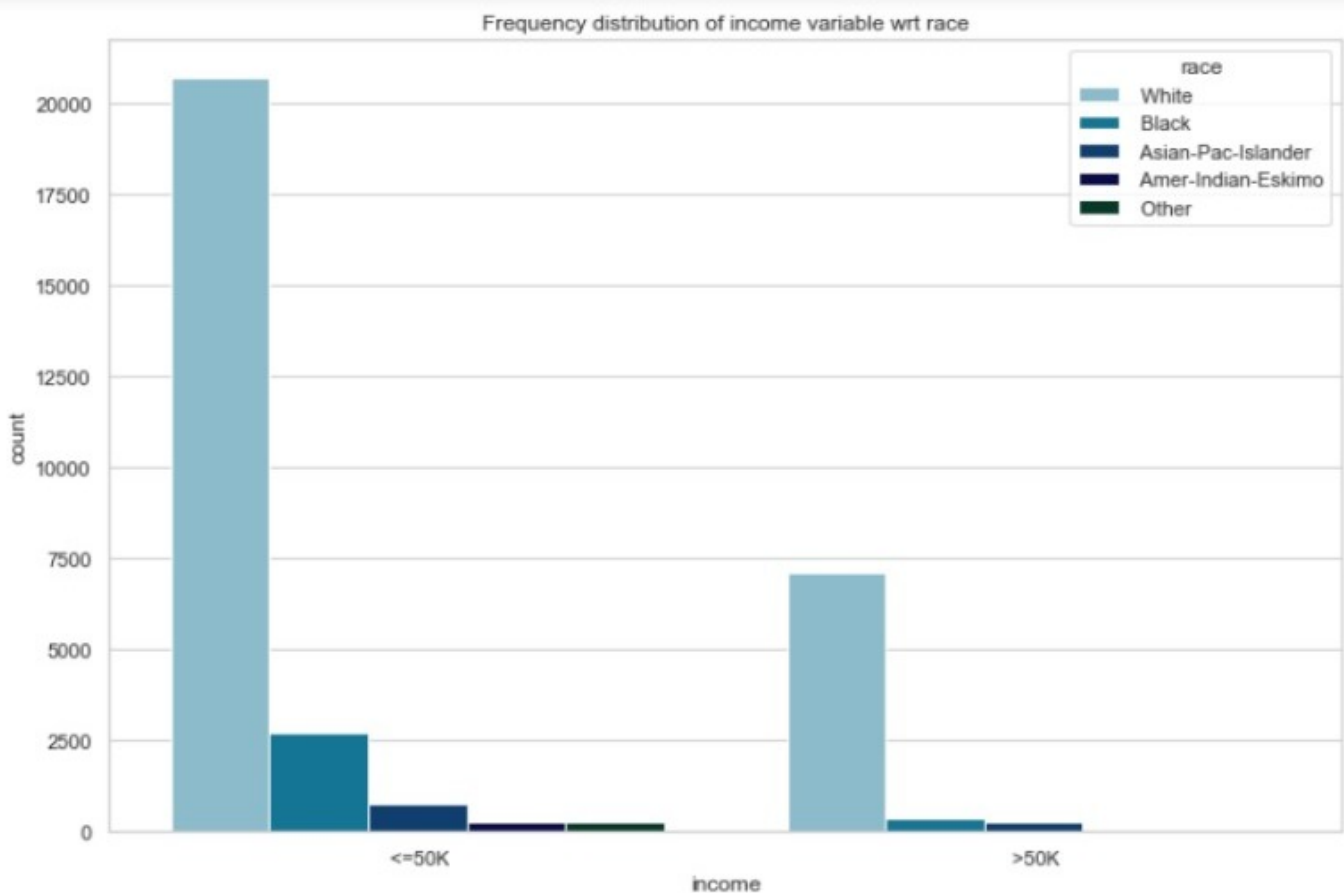
```
Out[53]: 9
```

```
In [54]: # view the unique labels
df.workclass.unique()
```

```
Out[54]: array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',
               ' Local-gov', ' ?', ' Self-emp-inc', ' Without-pay',
               ' Never-worked'], dtype=object)
```

```
In [55]: # view frequency distribution of values
df.workclass.value_counts()
```

```
Out[55]: Private      22696
Self-emp-not-inc    2541
Local-gov          2093
?                  1836
State-gov          1298
```





```
In [56]: # replace '?' values in workclass variable with `NaN`
```

```
df['workclass'].replace('?', np.NaN, inplace=True)
```

```
In [57]: # again check the frequency distribution of values in workclass variable
```

```
df.workclass.value_counts()
```

```
Out[57]: Private          22696  
Self-emp-not-inc      2541  
Local-gov            2093  
State-gov            1298  
Self-emp-inc         1116  
Federal-gov          960  
Without-pay          14  
Never-worked          7  
Name: workclass, dtype: int64
```

- Now, we can see that there are no values encoded as ? in the workclass variable.
- I will adopt similar approach with `occupation` and `native_country` column.

#### Visualize workclass variable

```
In [146]: f, ax = plt.subplots(figsize=(10, 6))  
ax = df.workclass.value_counts().plot(kind="bar", color="green")  
ax.set_title("Frequency distribution of workclass variable")  
ax.set_xticklabels(df.workclass.value_counts().index, rotation=30)
```

```
In [54]: # view the unique labels
```

```
df.workclass.unique()
```

```
Out[54]: array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',  
                ' Local-gov', ' ?', ' Self-emp-inc', ' Without-pay',  
                ' Never-worked'], dtype=object)
```

```
In [55]: # view frequency distribution of values
```

```
df.workclass.value_counts()
```

```
Out[55]: Private          22696  
Self-emp-not-inc      2541  
Local-gov            2093  
?                    1836  
State-gov            1298  
Self-emp-inc         1116  
Federal-gov          960  
Without-pay          14  
Never-worked          7  
Name: workclass, dtype: int64
```

We can see that there are 1836 values encoded as ? in workclass variable. I will replace these ? with NaN .

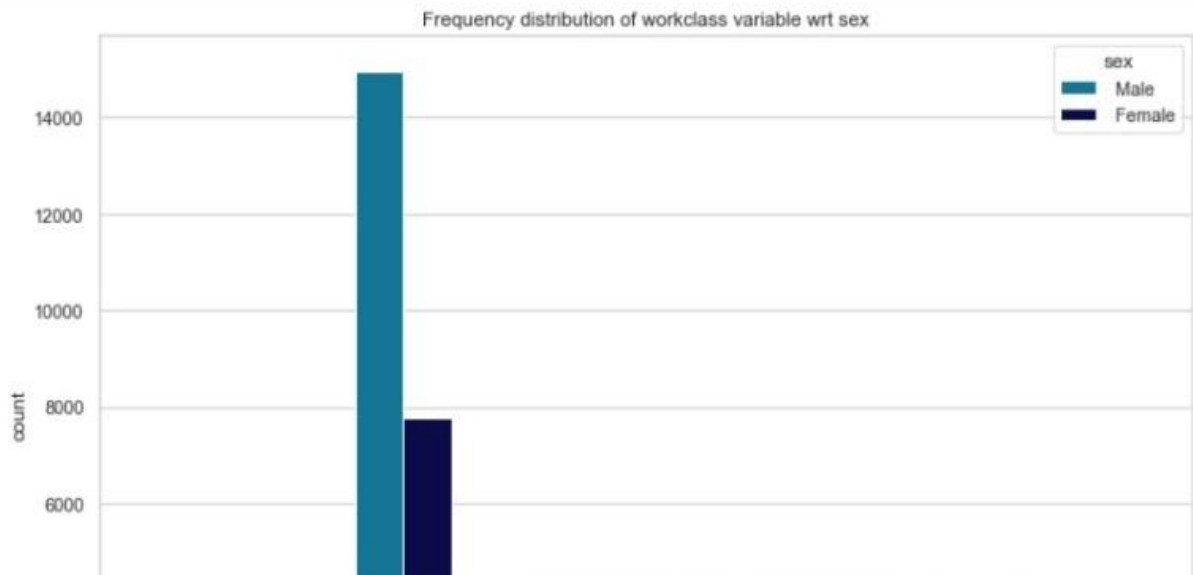
```
In [56]: # replace '?' values in workclass variable with `NaN`
```

```
df['workclass'].replace('?', np.NaN, inplace=True)
```

## Interpretation


- We can see that there are lot more private workers than other category of workers.

```
In [144]: categorical_features=["sex","income"]
for feature in categorical_features:
    f,ax = plt.subplots(figsize=(12, 8))
    ax = sns.countplot(x="workclass", hue=feature, data=df, palette="ocean_r")
    ax.set_title(f"Frequency distribution of workclass variable wrt {feature}")
    plt.show()
```




## Visualize workclass variable


```
In [146]: f, ax = plt.subplots(figsize=(10, 6))
ax = df.workclass.value_counts().plot(kind="bar", color="green")
ax.set_title("Frequency distribution of workclass variable")
ax.set_xticklabels(df.workclass.value_counts().index, rotation=30)
plt.show()
```



Frequency distribution of workclass variable wrt income



Frequency distribution of workclass variable wrt sex



Frequency distribution of workclass variable wrt income



```
In [63]: # view frequency distribution of values
```

```
df.occupation.value_counts()
```

```
Out[63]: Prof-specialty      4140
Craft-repair      4099
Exec-managerial    4066
Adm-clerical       3770
Sales             3650
Other-service      3295
Machine-op-inspct  2002
?                1843
Transport-moving   1597
Handlers-cleaners  1370
Farming-fishing    994
Tech-support       928
Protective-serv    649
Priv-house-serv    149
Armed-Forces        9
Name: occupation, dtype: int64
```

We can see that there are 1843 values encoded as `?` in occupation variable. I will replace these `?` with `NaN`.

```
In [64]: # replace '?' values in occupation variable with `NaN`
```

```
df['occupation'].replace('?', np.NaN, inplace=True)
```

```
In [65]: # again check the frequency distribution of values
```

## Interpretation

- We can see that workers make less than equal to 50k in most of the working categories.
- But this trend is more appealing in Private `workclass` category.
- We can see that there are more male workers than female workers in all the working category.
- The trend is more appealing in Private sector.

## Explore `occupation` variable

```
In [61]: # check number of unique labels
```

```
df.occupation.nunique()
```

```
Out[61]: 15
```

```
In [62]: # view unique labels
```

```
df.occupation.unique()
```

```
Out[62]: array([' Adm-clerical', ' Exec-managerial', ' Handlers-cleaners',
                ' Prof-specialty', ' Other-service', ' Sales', ' Craft-repair',
                ' Transport-moving', ' Farming-fishing', ' Machine-op-inspct',
                ' Tech-support', ' ?', ' Protective-serv', ' Armed-Forces',
                ' Priv-house-serv'], dtype=object)
```

```
In [65]: # again check the frequency distribution of values
```

```
df.occupation.value_counts()
```

```
Out[65]: Prof-specialty      4140
          Craft-repair      4099
          Exec-managerial   4066
          Adm-clerical      3770
          Sales             3650
          Other-service     3295
          Machine-op-inspct 2002
          Transport-moving  1597
          Handlers-cleaners 1370
          Farming-fishing   994
          Tech-support      928
          Protective-serv   649
          Priv-house-serv   149
          Armed-Forces      9
          Name: occupation, dtype: int64
```

```
In [147]: # visualize frequency distribution of `occupation` variable
```

```
f, ax = plt.subplots(figsize=(12, 8))
ax = sns.countplot(x="occupation", data=df, palette="ocean_r")
ax.set_title("Frequency distribution of occupation variable")
ax.set_xticklabels(df.occupation.value_counts().index, rotation=30)
plt.show()
```

```
In [69]: # check frequency distribution of values
```

```
df.native_country.value_counts()
```

```
Out[69]: United-States      29170
          Mexico           643
          ?                583
          Philippines      198
          Germany          137
          Canada           121
          Puerto-Rico      114
          El-Salvador      106
          India            100
          Cuba             95
          England          90
          Jamaica          81
          South            80
          China            75
          Italy            73
          Dominican-Republic 70
          Vietnam          67
          Guatemala        64
          Japan            62
          Poland           60
          Columbia         59
          Taiwan           51
          Haiti            44
          Iran             43
          Portugal         37
          Nicaragua        34
```

```
In [70]: # replace '?' values in native_country variable with `NaN`
```

```
df['native_country'].replace('?', np.NaN, inplace=True)
```

```
In [71]: # again check the frequency distribution of values
```

```
df.native_country.value_counts()
```

```
Out[71]: United-States      29170
Mexico                    643
Philippines              198
Germany                  137
Canada                   121
Puerto-Rico             114
El-Salvador              106
India                    100
Cuba                     95
England                  90
Jamaica                   81
South                    80
China                    75
Italy                    73
Dominican-Republic       70
Vietnam                   67
Guatemala                 64
Japan                     62
Poland                    60
Columbia                  59
Taiwan                    51
Haiti                     44
Iran                      43
```

```
Columbia                  59
Taiwan                    51
Haiti                     44
Iran                      43
Portugal                  37
Nicaragua                 34
Peru                      31
France                    29
Greece                    29
Ecuador                   28
Ireland                   24
Hong                      20
Cambodia                  19
Trinidad&Tobago           19
Laos                      18
Thailand                   18
Yugoslavia                16
Outlying-US(Guam-USVI-etc) 14
Honduras                  13
Hungary                   13
Scotland                  12
Holand-Netherlands        1
Name: native_country, dtype: int64
```

We can see that there are 583 values encoded as ? in native\_country variable. I will replace these ? with NaN .

```
In [70]: # replace '?' values in native_country variable with `NaN`
```

```
df['native_country'].replace('?', np.NaN, inplace=True)
```

Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Cambodia	19
Trinidad&Tobago	19
Laos	18
Thailand	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1


Name: native\_country, dtype: int64

In [148]: # visualize frequency distribution of `native\_country` variable

```
f, ax = plt.subplots(figsize=(16, 12))
ax = sns.countplot(x="native_country", data=df, palette="ocean_r")
ax.set_title("Frequency distribution of native_country variable")
ax.set_xticklabels(df.native_country.value_counts().index, rotation=90)
plt.show()
```

In [148]: # visualize frequency distribution of `native\_country` variable

```
f, ax = plt.subplots(figsize=(12, 12))
ax = sns.countplot(x="native_country", data=df, palette="ocean_r")
ax.set_title("Frequency distribution of native_country variable")
ax.set_xticklabels(df.native_country.value_counts().index, rotation=90)
plt.show()
```





We can see that `United-States` dominate amongst the `native_country` variables.

```
In [73]: #Checking missing categorical values
df[categorical].isnull().sum()
```

```
Out[73]: workclass      1836
education      0
marital_status 0
occupation     1843
relationship    0
race            0
sex            0
native_country  583
income         0
dtype: int64
```

Now, we can see that `workclass`, `occupation` and `native_country` variable contains missing values.

## 6. Explore Numerical Variables

```
In [75]: numerical = [var for var in df.columns if df[var].dtype!='O']

print('There are {} numerical variables\n'.format(len(numerical)))

print('The numerical variables are :\n\n', numerical)

There are 6 numerical variables
```

```
In [75]: numerical = [var for var in df.columns if df[var].dtype!='O']

print('There are {} numerical variables\n'.format(len(numerical)))

print('The numerical variables are :\n\n', numerical)
```

There are 6 numerical variables

The numerical variables are :

`['age', 'fnlwgt', 'education_num', 'capital_gain', 'capital_loss', 'hours_per_week']`

```
In [76]: df[numerical].head()
```

```
Out[76]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
0	39	77516	13	2174	0	40
1	50	83311	13	0	0	13
2	38	215646	9	0	0	40
3	53	234721	7	0	0	40
4	28	338409	13	0	0	40

```
In [77]: df[numerical].isnull().sum()
```

```
Out[77]: age            0
fnlwgt            0
education_num     0
capital_gain      0
```

```
In [77]: df[numerical].isnull().sum()
```

```
Out[77]: age                0
         fnlwgt             0
         education_num      0
         capital_gain       0
         capital_loss       0
         hours_per_week     0
         dtype: int64
```

We can see that there are no missing values in the numerical variables.

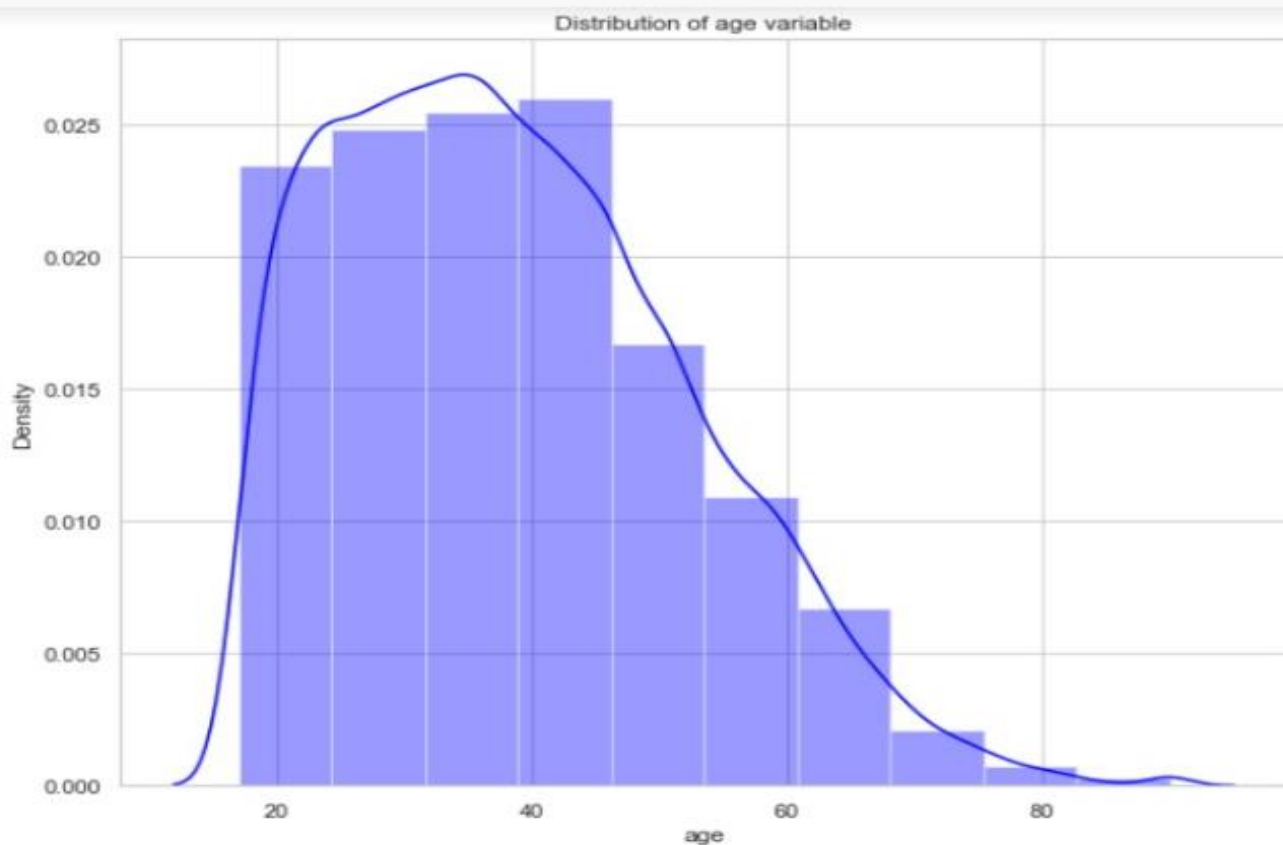
### Explore age variable

```
In [78]: df['age'].nunique()
```

```
Out[78]: 73
```


### View the distribution of age variable

```
In [79]: f, ax = plt.subplots(figsize=(10,8))
         x = df['age']
         ax = sns.distplot(x, bins=10, color='blue')
         ax.set_title("Distribution of age variable")
         plt.show()
```



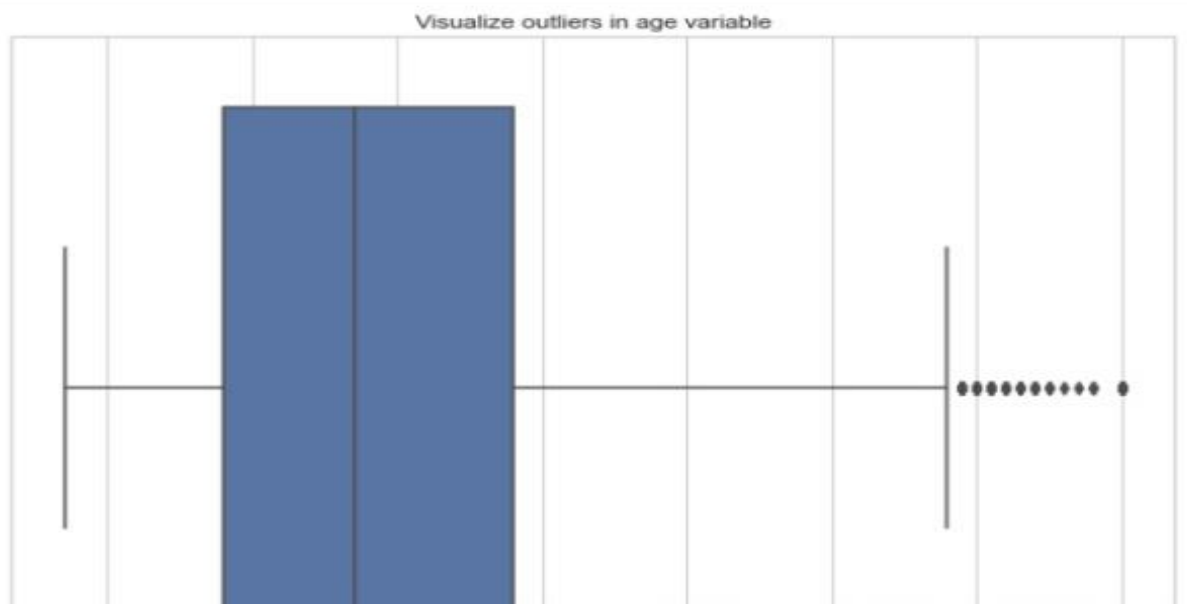
We can shade under the density curve and use a different color as follows:-

```
In [81]: f, ax = plt.subplots(figsize=(10,8))
x = df['age']
x = pd.Series(x, name="Age variable")
ax = sns.kdeplot(x, shade=True, color='red')
ax.set_title("Distribution of age variable")
plt.show()
```



**Detect outliers in age variable with boxplot**

```
In [82]: f, ax = plt.subplots(figsize=(10,8))
x = df['age']
ax = sns.boxplot(x)
ax.set_title("Visualize outliers in age variable")
plt.show()
```




We can see that there are lots of outliers in `age` variable.

### Explore relationship between `age` and `income` variables

```
In [149]: f, ax = plt.subplots(figsize=(10, 8))
ax = sns.boxplot(x="income", y="age", data=df)
ax.set_title("Visualize income wrt age variable")
plt.show()

f, ax = plt.subplots(figsize=(10, 8))
ax = sns.boxplot(x="income", y="age", data=df)
ax.set_title("Visualize income wrt age variable")
plt.show()
```






## Interpretation

- As expected, younger people make less money as compared to senior people.
- Whites are more older than other groups of people.

## Find out the correlations

```
In [137]: # plot correlation heatmap to find out correlations
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(),annot=True,cmap="ocean_r")
```

Out[137]: <AxesSubplot:>

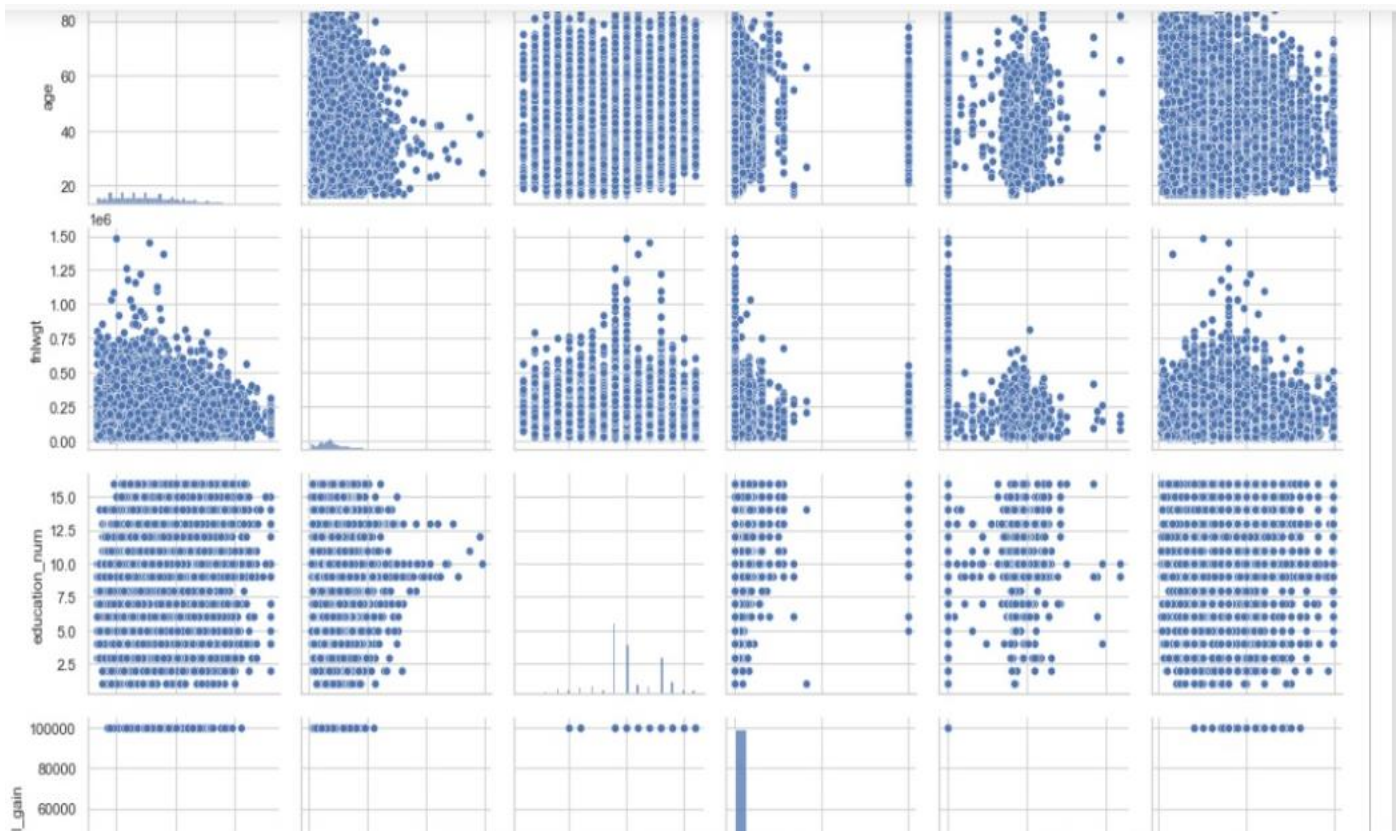
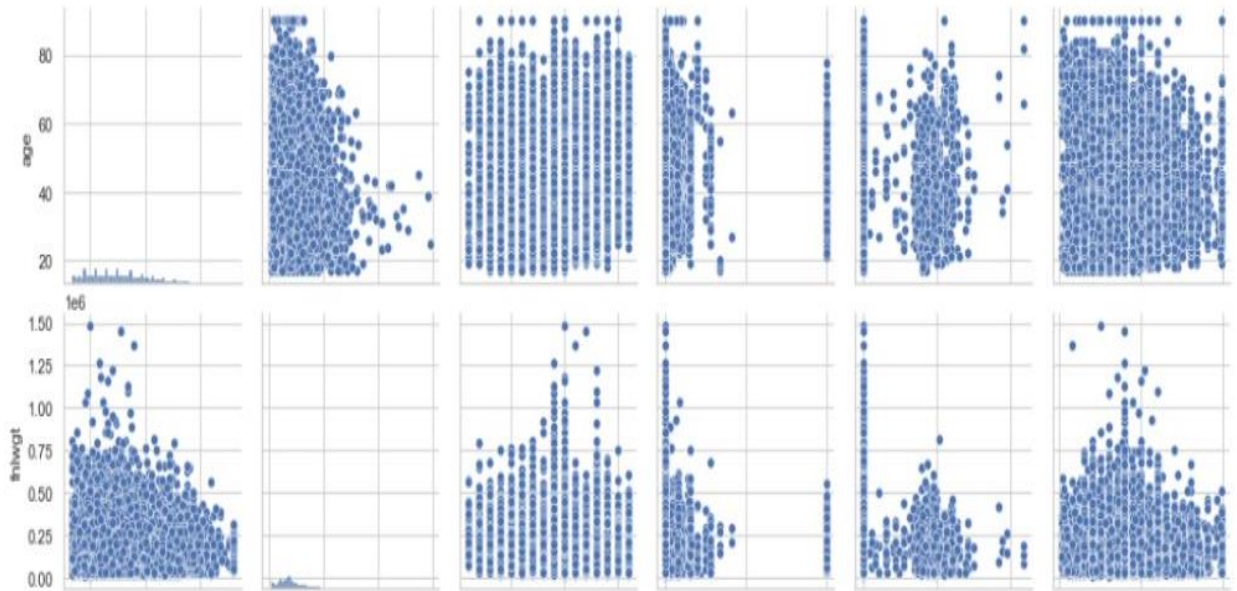



## Interpretation

- We can see that there is no strong correlation between variables.

## Plot pairwise relationships in dataset

```
In [88]: sns.pairplot(df)
plt.show()
```

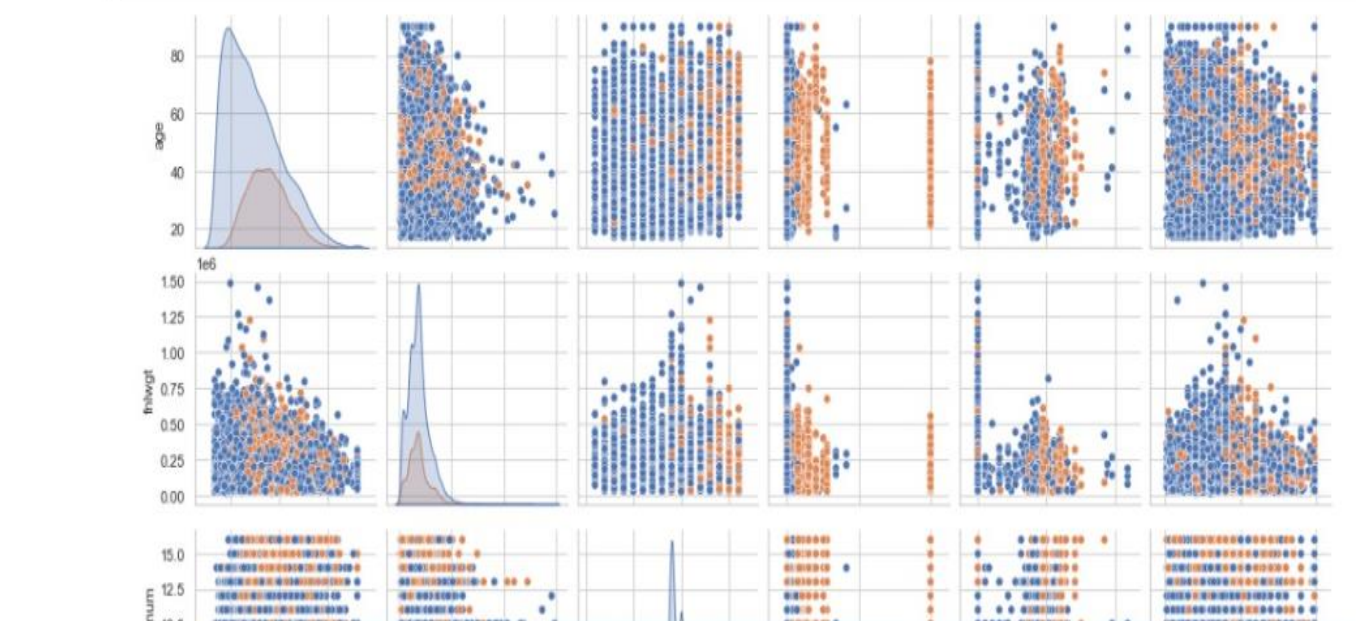




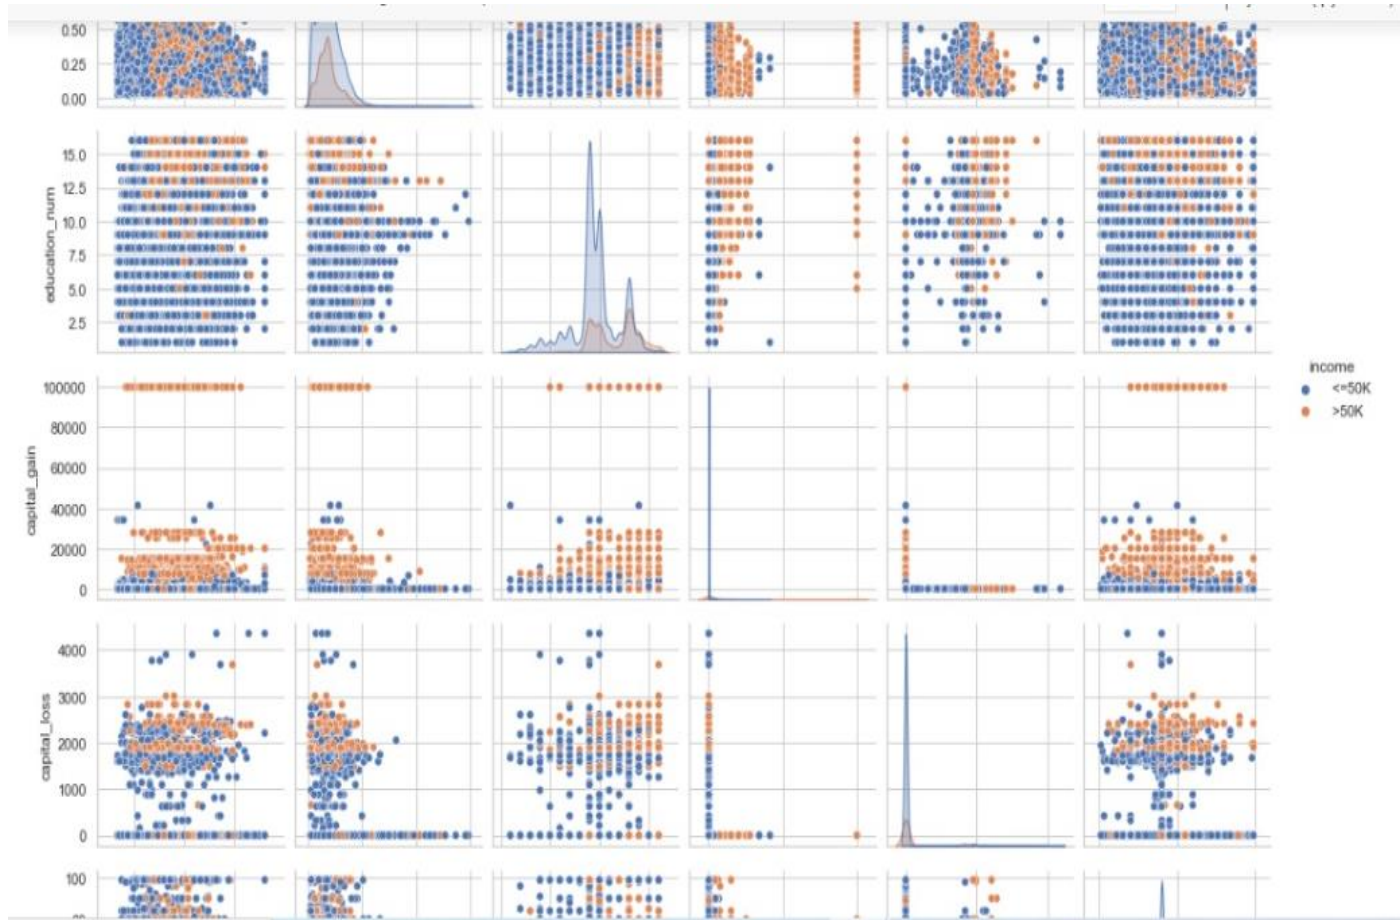
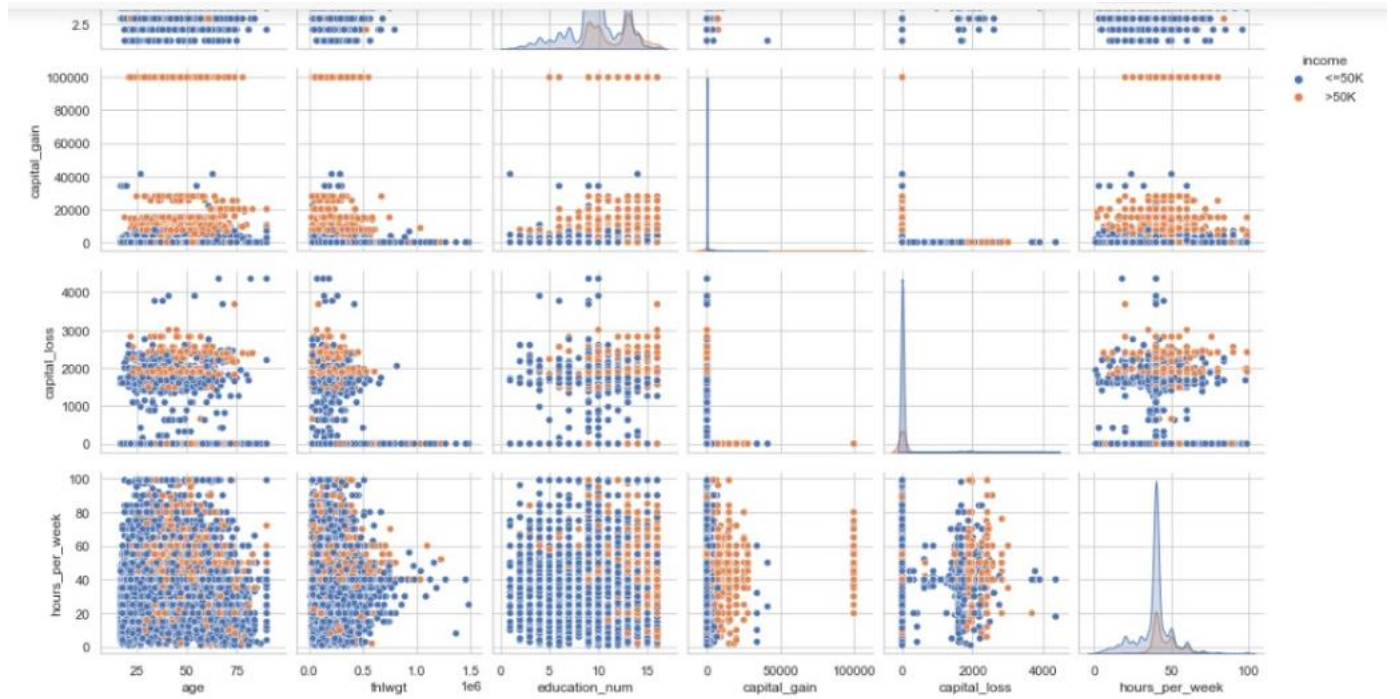
### Interpretation

- We can see that `age` and `fnlwtgt` are positively skewed.
- The variable `education_num` is negatively skewed while `hours_per_week` is normally distributed.
- There exists weak positive correlation between `capital_gain` and `education_num` (correlation coefficient=0.1226).

```
In [89]: sns.pairplot(df, hue="income")
plt.show()
```







## 7. Declare feature vector and target variable

```
In [91]: X = df.drop(['income'], axis=1)
        y = df['income']
```

## 8. Split data into separate training and test set

```
In [92]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
```

```
In [93]: # check the shape of X_train and X_test
        X_train.shape, X_test.shape
```

```
Out[93]: ((22792, 14), (9769, 14))
```

## 9. Feature Engineering .

### 9.1 Display categorical variables in training set

## 9. Feature Engineering .

### 9.1 Display categorical variables in training set

```
In [94]: categorical = [col for col in X_train.columns if X_train[col].dtypes == 'O']
        categorical
```

```
Out[94]: ['workclass',
          'education',
          'marital_status',
          'occupation',
          'relationship',
          'race',
          'sex',
          'native_country']
```

### 9.2 Display numerical variables in training set

```
In [95]: numerical = [col for col in X_train.columns if X_train[col].dtypes != 'O']
        numerical
```

```
Out[95]: ['age',
          'fnlwgt',
          'education_num',
          'capital_gain',
          'capital_loss',
```

### 9.3 Engineering missing values in categorical variables

In [96]: *# print percentage of missing values in the categorical variables in training set*

```
X_train[categorical].isnull().mean()
```

Out[96]:

workclass	0.055985
education	0.000000
marital_status	0.000000
occupation	0.056072
relationship	0.000000
race	0.000000
sex	0.000000
native_country	0.018164

dtype: float64

In [97]: *# print categorical variables with missing data*

```
for col in categorical:
    if X_train[col].isnull().mean()>0:
        print(col, (X_train[col].isnull().mean()))
```

workclass 0.055984555984555984  
occupation 0.05607230607230607  
native\_country 0.018164268164268166

In [98]: *# impute missing categorical variables with most frequent value*

```
for df2 in [X_train, X_test]:
    df2['workclass'].fillna(X_train['workclass'].mode()[0], inplace=True)
    df2['occupation'].fillna(X_train['occupation'].mode()[0], inplace=True)
    df2['native_country'].fillna(X_train['native_country'].mode()[0], inplace=True)
```

In [99]: *# check missing values in categorical variables in X\_train*

```
X_train[categorical].isnull().sum()
```

Out[99]:

workclass	0
education	0
marital_status	0
occupation	0
relationship	0
race	0
sex	0
native_country	0

dtype: int64

In [100]: *# check missing values in categorical variables in X\_test*

```
X_test[categorical].isnull().sum()
```

Out[100]:

workclass	0
education	0
marital_status	0
occupation	0



```
In [101]: # check missing values in X_train
X_train.isnull().sum()
```

```
Out[101]: age                0
workclass            0
fnlwgt              0
education            0
education_num        0
marital_status       0
occupation           0
relationship         0
race                0
sex                 0
capital_gain         0
capital_loss         0
hours_per_week       0
native_country       0
dtype: int64
```

```
In [102]: # check missing values in X_test
X_test.isnull().sum()
```

```
Out[102]: age                0
workclass            0
fnlwgt              0
education            0
education_num        0
marital_status       0
```

```
occupation           0
relationship         0
race                0
sex                 0
capital_gain         0
capital_loss         0
hours_per_week       0
native_country       0
dtype: int64
```

We can see that there are no missing values in X\_train and X\_test.

## 9.4 Encode categorical variables

```
In [103]: # preview categorical variables in X_train
```

```
X_train[categorical].head()
```

Out[103]:

	workclass	education	marital_status	occupation	relationship	race	sex	native_country
32098	Private	HS-grad	Married-civ-spouse	Craft-repair	Husband	White	Male	United-States
25206	State-gov	HS-grad	Divorced	Adm-clerical	Unmarried	White	Female	United-States
23491	Private	Some-college	Married-civ-spouse	Sales	Husband	White	Male	United-States
12367	Private	HS-grad	Never-married	Craft-repair	Not-in-family	White	Male	Guatemala
7054	Private	7th-8th	Never-married	Craft-repair	Not-in-family	White	Male	Germany

```
In [104]: # import category encoders
```

```
import category_encoders as ce
```

```
In [105]: # encode categorical variables with one-hot encoding
```

```
encoder = ce.OneHotEncoder(cols=['workclass', 'education', 'marital_status', 'occupation', 'relationship',  
                                'race', 'sex', 'native_country'])
```

```
X_train = encoder.fit_transform(X_train)
```

```
X_test = encoder.transform(X_test)
```

```
In [106]: X_train.head()
```

Out[106]:

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7	workclass_8	fnlwgt	...	native_country_32	native_c
32098	45	1	0	0	0	0	0	0	0	170871	...	0	
25206	47	0	1	0	0	0	0	0	0	108890	...	0	
23491	48	1	0	0	0	0	0	0	0	187505	...	0	
12367	29	1	0	0	0	0	0	0	0	145592	...	0	
7054	23	1	0	0	0	0	0	0	0	203003	...	0	

5 rows × 105 columns



```
In [107]: X_train.shape
```



```
In [109]: X_test.shape
```

```
Out[109]: (9769, 105)
```

- We now have training and testing set ready for model building. Before that, we should map all the feature variables onto the same scale. It is called **feature scaling**. We will do it as follows.

## 10. Feature Scaling

```
In [110]: cols = X_train.columns
```

```
In [111]: from sklearn.preprocessing import RobustScaler
```

```
scaler = RobustScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
In [112]: X_train = pd.DataFrame(X_train, columns=cols)
```

```
In [113]: X_test = pd.DataFrame(X_test, columns=cols)
```

We now have X\_train dataset ready to be fed into the Random Forest classifier. We will do it as follows.

```
In [107]: X_train.shape
```

```
Out[107]: (22792, 105)
```

We can see that from the initial 14 columns, we now have 105 columns in training set.

Similarly, I will take a look at the X\_test set.

```
In [108]: X_test.head()
```

```
Out[108]:
```

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7	workclass_8	fnlwgt	...	native_country_32	native_c
22278	27	1	0	0	0	0	0	0	0	177119	...	0	0
8950	27	1	0	0	0	0	0	0	0	216481	...	0	0
7838	25	1	0	0	0	0	0	0	0	256263	...	0	0
16505	46	1	0	0	0	0	0	0	0	147640	...	0	0
19140	45	1	0	0	0	0	0	0	0	172822	...	0	0

5 rows × 105 columns



```
In [109]: X_test.shape
```

```
Out[109]: (9769, 105)
```

We now have X\_train dataset ready to be fed into the Random Forest classifier. We will do it as follows.

## 11. Random Forest Classifier model with default parameters

```
In [114]: # import Random Forest classifier
          from sklearn.ensemble import RandomForestClassifier

          # instantiate the classifier
          rfc = RandomForestClassifier(random_state=0)

          # fit the model
          rfc.fit(X_train, y_train)

          # Predict the Test set results
          y_pred = rfc.predict(X_test)
```

Here, I have build the Random Forest Classifier model with default parameter of `n_estimators = 10`. So, I have used 10 decision-trees to build the model. Now, I will increase the number of decision-trees and see its effect on accuracy.

## 12. Random Forest Classifier model with 100 Decision Trees)

```
In [115]: # instantiate the classifier with n_estimators = 100
          rfc_100 = RandomForestClassifier(n_estimators=100, random_state=0)

          # fit the model to the training set
          rfc_100.fit(X_train, y_train)

          # Predict on the test set results
          y_pred_100 = rfc_100.predict(X_test)

          # Check accuracy score
          print('Model accuracy score with 100 decision-trees : {0:0.4f}'.format(accuracy_score(y_test, y_pred_100)))
```

The model accuracy score with 10 decision-trees is 0.8446 but the same with 100 decision-trees is 0.8521. So, as expected accuracy increases with number of decision-trees in the model.

### 13. Find important features with Random Forest model

Until now, We have used all the features given in the model. Now, I will select only the important features, build the model using these features and see its effect on accuracy.

First, We will create the Random Forest model as follows:-

```
In [116]: # create the classifier with n_estimators = 100

clf = RandomForestClassifier(n_estimators=100, random_state=0)

# fit the model to the training set

clf.fit(X_train, y_train)
```

```
Out[116]: RandomForestClassifier(random_state=0)
```

Now, We will use the feature importance variable to see feature importance scores.

```
In [117]: # view the feature scores
```

Now, I will build the random forest model again and check accuracy.

```
In [120]: # instantiate the classifier with n_estimators = 100

clf = RandomForestClassifier(n_estimators=100, random_state=0)

# fit the model to the training set

clf.fit(X_train, y_train)

# Predict on the test set results

y_pred = clf.predict(X_test)

# Check accuracy score

print('Model accuracy score with native_country_41 variable removed : {0:0.4f}'.format(accuracy_score(y_test, y_pred)))
```

Model accuracy score with native\_country\_41 variable removed : 0.8544

#### Interpretation

- I have removed the native\_country\_41 variable from the model, rebuild it and checked its accuracy.

```
Out[117]: fnlwgt      0.159772
age          0.149074
capital_gain 0.091299
hours_per_week 0.086339
education_num 0.065130
...
native_country_16 0.000028
occupation_14     0.000015
native_country_35 0.000009
workclass_8       0.000008
native_country_41 0.000000
Length: 105, dtype: float64
```

We can see that the most important feature is `fnlwgt` and least important feature is `native_country_41`.

## 14. Build the Random Forest model on selected features

Now, We will drop the least important feature `native_country_41` from the model, rebuild the model and check its effect on accuracy.

```
In [119]: # drop the least important feature from X_train and X_test

X_train = X_train.drop(['native_country_41'], axis=1)
X_test = X_test.drop(['native_country_41'], axis=1)
```

Now, I will build the random forest model again and check accuracy.

### Interpretation

- I have removed the `native_country_41` variable from the model, rebuild it and checked its accuracy.
- The accuracy of the model now comes out to be 0.8544.
- The accuracy of the model with all the variables taken into account is 0.8521.
- So, we can see that the model accuracy has been improved with `native_country_41` variable removed from the model.

Now, based on the above analysis we can conclude that our classification model accuracy is very good. Our model is doing a very good job in terms of predicting the class labels.

But, it does not give the underlying distribution of values. Also, it does not tell anything about the type of errors our classifier is making.

We have another tool called `Confusion matrix` that comes to our rescue.

```
In [121]: # Print the Confusion Matrix and slice it into four pieces

from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)

print('Confusion matrix\n\n', cm)
```

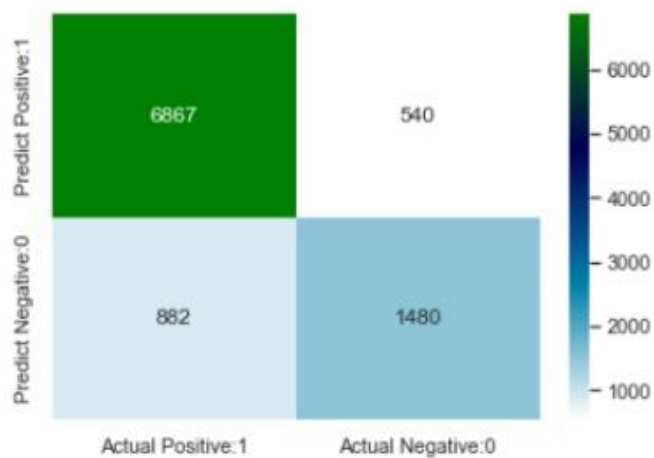
Confusion matrix

```
[[6867  540]
 [ 882 1480]]
```

```
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                          index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='ocean_r')
```

Out[139]: <AxesSubplot:>



## 15. Classification Report

```
In [123]: from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support
```

## 15. Classification Report

```
In [123]: from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))
```

```
              precision    recall  f1-score   support

<=50K         0.89         0.93         0.91         7407
>50K          0.73         0.63         0.68         2362

accuracy              0.85         0.85         0.85         9769
macro avg         0.81         0.78         0.79         9769
weighted avg         0.85         0.85         0.85         9769
```



## 17. Results and Conclusion

1. In this project, We build a Random Forest Classifier to predict the income of a person. We build two models, one with 10 decision-trees and another one with 100 decision-trees.
2. The model accuracy score with 10 decision-trees is 0.8446 but the same with 100 decision-trees is 0.8521 . So, as expected accuracy increases with number of decision-trees in the model.
3. We have used the Random Forest model to find only the important features, build the model using these features and see its effect on accuracy.
4. We have removed the native\_country\_41 variable from the model, rebuild it and checked its accuracy. The accuracy of the model with native\_country\_41 variable removed is 0.8544 . So, we can see that the model accuracy has been improved with native\_country\_41 variable removed from the model.
5. Confusion matrix and classification report are another tool to visualize the model performance. They yield good performance.

That is the end of this kernel. We hope you find it useful and enjoyable.

Your feedback and comments are most welcome.

### THANKS TO ALL THE GROUP MEMBERS: =

- a. Riddhey vyas
- b. Deeptimayee Maharana
- c. 483 chaithanya
- d. Akhil Kumar reddy
- e. Sumanth polieni
- f. Thota samanvitha
- g. Neeha shaik
- h. Usha Sri
- i. Ranhith T

*Thank you*